

# Measuring the Cost of Volatile in Java

Stefan Nüesch  
Seminar Software Composition 2013

# Overview

- Motivation
- Volatile Semantics
- Micro-Benchmark
- Displaying the data
- Results
- Questions

# The Problem

- Performance impact of volatile variables often not clear
- -> How do resource contention, read / write ratio, architecture and spatial locality in memory influence the performance of accessing volatile variables in Java?

# Scope

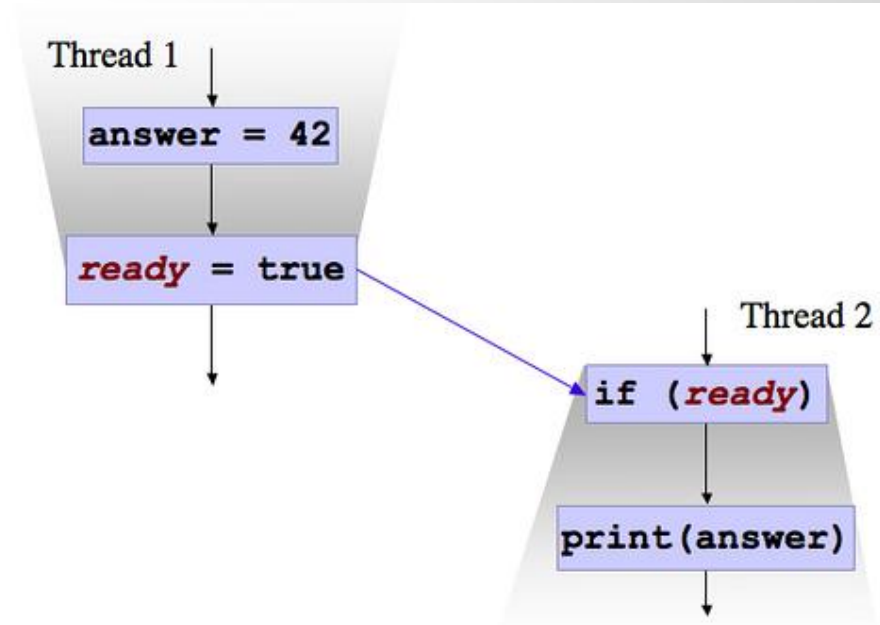
- In scope: compare performance for different levels of contention, different read / write ratios and other parameters
- Out of scope: Comparing `volatile` to `synchronized`, locking or other constructs

# Methodology


- Varying parameter settings systematically
- Warm-up runs
- Multiple run cycles
- Take averages for comparison
- Same number of accesses per run

# Volatile Semantics (Java 1.5+)

- volatile writes are always written through to memory
- **volatile** established a "happens before" relationship
- reads / writes subsequent to a write cannot be reordered



# Micro-Benchmark



# Micro-Benchmark (cont'd)

- Parameters:
  - Number of Keys
  - Key length
  - Number of Threads
  - Number of read / write cycles
  - Read / write ratio
  - Keyset overlapping
  - Number of run cycles



# Plotting Results


- Python script for collecting the data and preparing the plot
- gnuplot (with python-gnuplot) used for drawing the plot
- single plots created automatically

# Plotting Results - Python Code


```
plot = Gnuplot.Gnuplot()
plot.title(set)
plot('set term pdf font "Helvetica, 10"')
plot('set style line 1 lt 1 lc rgb "#FF0000" lw 3 # red')
plot('set output "'+resultDir+set.replace(' ', '_')+'.pdf";')
plot('set ylabel "'+yAxisLabel+'";')
plot('set xlabel "'+xAxisLabel[varPosition]+'";')
plot('set key top left;')
plot('set yrange [:]')
plot('set xrange [:]')
plot('plot "'+resultDir+dataFile+'" using ($1):($2*0.000001):($3*0.000001):($4*0.000001)
with yerrorlines title "'+xAxisLabel[varPosition]+'";')
```

# Results


3 byte keys, 17 Quad Core




10240 byte keys, 17 Quad Core



3 byte keys, AMD 64 Dual Core



102400 byte keys, AMD 64 Dual Core



# Results

- No impact of read contention
- Considerable performance impact with little writing -> factor 3 on AMD 4
- Performance of writing is strongly affected by contention -> factor 3



# Future Work

- Assess impact of memory layout
- Compare several architectures
- In depth analysis of certain behaviour